

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Dataset

<https://www.kaggle.com/datasets/rashikrahmanpritom/heart-attack-analysis-prediction-dataset>

```
df=pd.read_csv("heart.csv")
df.head()
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
# Predictor variables (X)
X =df.drop(columns=['output'])

# Response variable (Y)
Y = df['output']



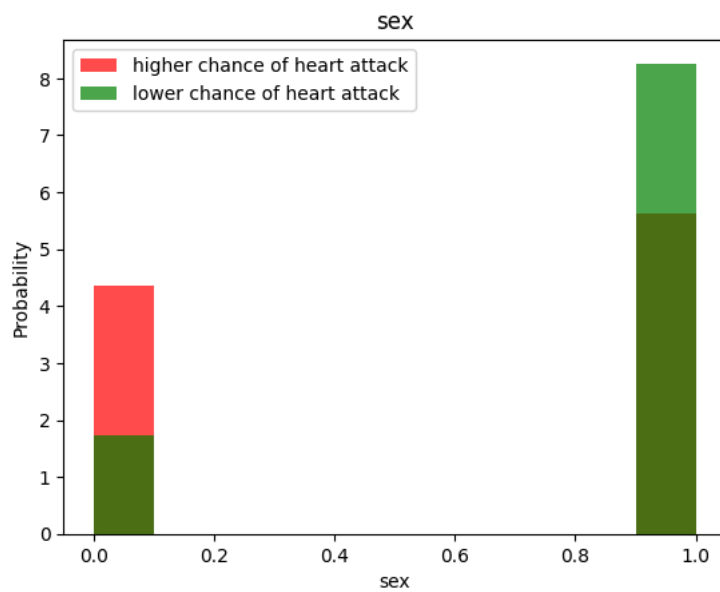
# Print the predictor and response variable
print("Predictor variables (X):\n", X.head())
print("\nResponse variable (Y):\n", Y.head())
```

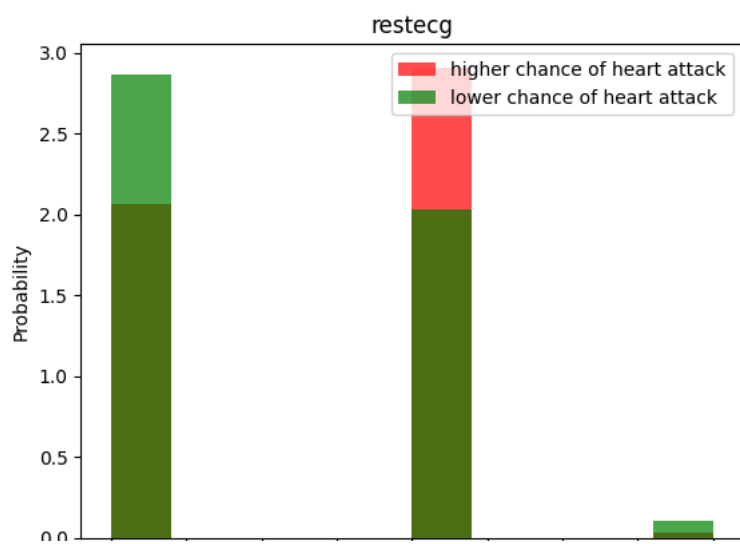
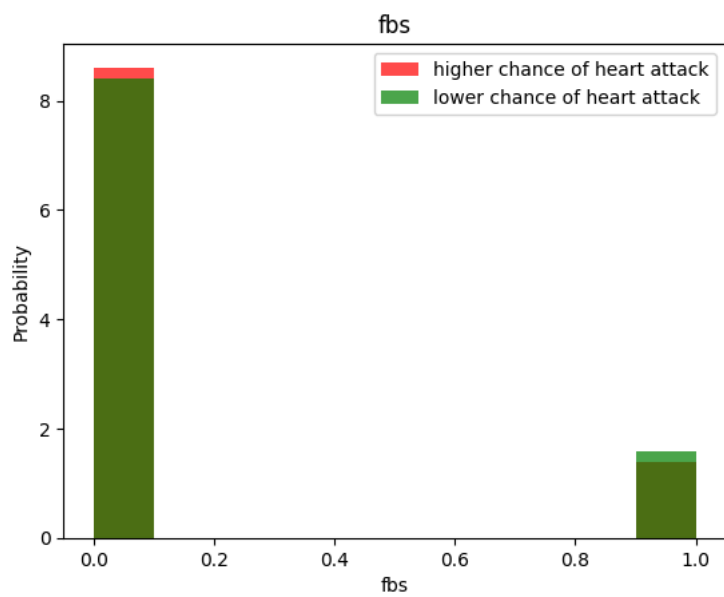
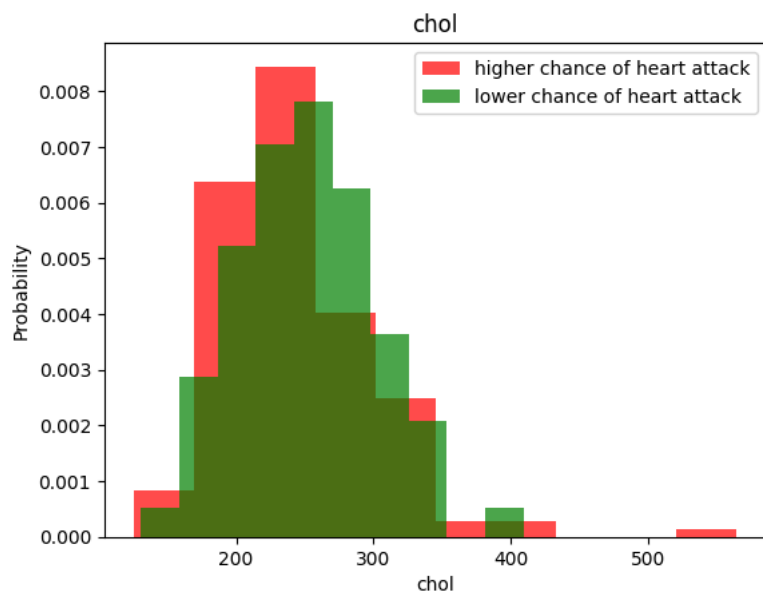
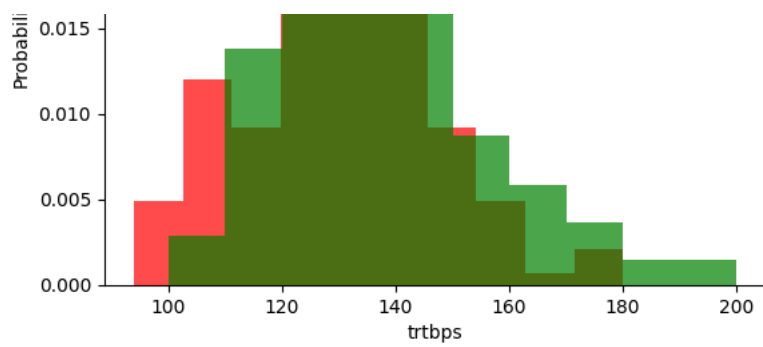
```
Predictor variables (X):
   age  sex  cp  trtbps  chol  fbs  restecg  thalachh  exng  oldpeak  slp  \
0   63   1   3   145   233   1         0     150     0     2.3     0
1   37   1   2   130   250   0         1     187     0     3.5     0
2   41   0   1   130   204   0         0     172     0     1.4     2
3   56   1   1   120   236   0         1     178     0     0.8     2
4   57   0   0   120   354   0         1     163     1     0.6     2




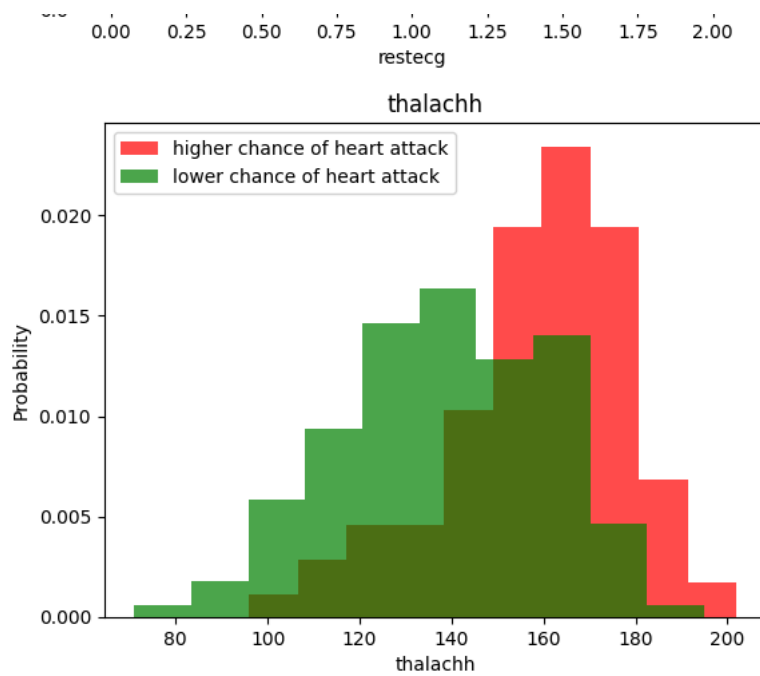
   caa  thall
0     0     1
1     0     2
2     0     2
3     0     2
4     0     2
```

```
Response variable (Y):
0     1
1     1
2     1
3     1
4     1
Name: output, dtype: int64
```

```
for label in df.columns[:-1]:
    plt.hist(df[df["output"]==1][label],color='red',label="higher chance of heart attack",alpha=0.7,density=True)
    plt.hist(df[df["output"]==0][label],color='green',label="lower chance of heart attack",alpha=0.7,density=True)
    plt.title(label)
    plt.ylabel("Probability")
    plt.xlabel(label)
    plt.legend()
    plt.show()
```





Splitting the data into train & test data 80:20 (First Case)

```
train,test=np.split(df.sample(frac=1),[int(0.8*len(df))])
```

```
# Importing the required modules for scaling data
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import SMOTE

def scale_dataset(data, oversample=False):

    # Scale the data
    scaler = StandardScaler()
    scaled_data = scaler.fit_transform(data.drop('output', axis=1))

    # Separate features and labels
    features = scaled_data
    labels = data['output']

    # Oversample the data
    if oversample:
        smote = SMOTE(random_state=42)
        features, labels = smote.fit_resample(features, labels)

    return data, features, labels

# Scale the train and test data
train, X_train, Y_train = scale_dataset(train, oversample=True)
test, X_test, Y_test = scale_dataset(test, oversample=True)
```

```
# Print shapes of the arrays
print("Shapes for train, X_train, and Y_train:")
print("Train shape:", train.shape)
print("X_train shape:", X_train.shape)
print("Y_train shape:", Y_train.shape)

print("\nShapes for test, X_test, and Y_test:")
print("Test shape:", test.shape)
print("X_test shape:", X_test.shape)
print("Y_test shape:", Y_test.shape)
```

```
Shapes for train, X_train, and Y_train:
Train shape: (242, 14)
X_train shape: (270, 13)
Y_train shape: (270,)

Shapes for test, X_test, and Y_test:
Test shape: (61, 14)
X_test shape: (62, 13)
Y_test shape: (62,)
```

Applying different classifications models on data (KNN,SVM,Naivebayes,Logistic Regression and LDA classifier) [80:20]

```
#Will give confusion matrix and classification report for each classifier for two thresholds 0.5 & 0.6 for 80:20 train-test.
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.metrics import confusion_matrix, classification_report

# A function to train and evaluate a classifier
def train_and_evaluate(classifier, X_train, Y_train, X_test, Y_test,threshold=0.5):
    # Train the classifier
    classifier.fit(X_train, Y_train)

    # Predict probabilities
    Y_prob = classifier.predict_proba(X_test)[:, 1]

    # Apply threshold
    Y_pred = (Y_prob > threshold).astype(int)

    ts
```

```
# Print confusion matrix
print("Confusion Matrix (Threshold = {}):".format(threshold))
print(confusion_matrix(Y_test, Y_pred))

# Print classification report
print("\nClassification Report (Threshold = {}):".format(threshold))
print(classification_report(Y_test, Y_pred))

# Initialize classifiers
knn_classifier = KNeighborsClassifier(n_neighbors=3)
svm_classifier = SVC(probability=True) # Set probability=True for SVC to enable probability estimates
nb_classifier = GaussianNB()
lr_classifier = LogisticRegression()
lda_classifier = LinearDiscriminantAnalysis()

# Train and evaluate classifiers
classifiers = [knn_classifier, svm_classifier, nb_classifier, lr_classifier, lda_classifier]
classifier_names = ['KNN', 'SVM', 'Naive Bayes', 'Logistic Regression', 'LDA']

# Define threshold values
thresholds = [0.5, 0.6]

# Loop over classifiers and threshold values
for classifier, name in zip(classifiers, classifier_names):
    for threshold in thresholds:
        print("\nTraining and evaluating", name, "classifier with threshold", threshold)
        train_and_evaluate(classifier, X_train, Y_train, X_test, Y_test, threshold)
```

```

from sklearn.metrics import roc_curve, auc

# Define a function to plot ROC-AUC curve
def plot_roc_curve(classifier, X_test, Y_test, name):
    # Predict probabilities
    Y_prob = classifier.predict_proba(X_test)[:, 1]

    # Calculate false positive rate, true positive rate, and thresholds
    fpr, tpr, thresholds = roc_curve(Y_test, Y_prob)


    # Calculate area under the curve (AUC)
    roc_auc = auc(fpr, tpr)

    # Plot ROC curve
    plt.plot(fpr, tpr, label=name + ' (AUC = %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], linestyle='--', color='gray', label='Random Guess')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve')
    plt.legend()
    plt.grid(True)

for classifier, name in zip(classifiers, classifier_names):
    classifier.fit(X_train, Y_train)

# Plot ROC curves for all classifiers
plt.figure(figsize=(10, 6))
for classifier, name in zip(classifiers, classifier_names):
    plot_roc_curve(classifier, X_test, Y_test, name)
plt.legend(loc='lower right')
plt.show()

```



✓ Splitting the data into train & test data 70:30(Second Case)

```
train2, test2 = np.split(df.sample(frac=1), [int(0.7*len(df))])
```

```

# Scale the train and test data
train2, X_train2, Y_train2 = scale_dataset(train2, oversample=True)
test2, X_test2, Y_test2 = scale_dataset(test2, oversample=True)

```

```

# Print shapes of the arrays
print("Shapes for train, X_train, and Y_train:")

```



```

print("Train shape:", train2.shape)
print("X_train shape:", X_train2.shape)
print("Y_train shape:", Y_train2.shape)

print("\nShapes for test, X_test, and Y_test:")
print("Test shape:", test2.shape)
print("X_test shape:", X_test2.shape)
print("Y_test shape:", Y_test2.shape)

```

```

Shapes for train, X_train, and Y_train:
Train shape: (212, 14)
X_train shape: (220, 13)
Y_train shape: (220,)

```

```

Shapes for test, X_test, and Y_test:
Test shape: (91, 14)
X_test shape: (110, 13)
Y_test shape: (110,)

```

Applying different classifications models on data (KNN,SVM,Naivebayes,Logestic Regression and LDA classifier) [70:30]

#Will give confusion matrix and classification report for each classifier for two thresholds 0.5 & 0.6 for 70:30 train-test.

```

# Loop over classifiers and threshold values
for classifier, name in zip(classifiers, classifier_names):
    for threshold in thresholds:
        print("\nTraining and evaluating", name, "classifier with threshold", threshold)
        train_and_evaluate(classifier, X_train2, Y_train2, X_test2, Y_test2, threshold)

```

```

Training and evaluating KNN classifier with threshold 0.5
Confusion Matrix (Threshold = 0.5):
[[41 14]
 [10 45]]

```

```

Classification Report (Threshold = 0.5):
              precision    recall  f1-score   support

     0       0.80      0.75      0.77        55
     1       0.76      0.82      0.79        55

 accuracy          0.78
macro avg          0.78      0.78      0.78        110
weighted avg       0.78      0.78      0.78        110

```

```

Training and evaluating KNN classifier with threshold 0.6
Confusion Matrix (Threshold = 0.6):
[[41 14]
 [10 45]]

```

```

Classification Report (Threshold = 0.6):
              precision    recall  f1-score   support

     0       0.80      0.75      0.77        55
     1       0.76      0.82      0.79        55

 accuracy          0.78
macro avg          0.78      0.78      0.78        110
weighted avg       0.78      0.78      0.78        110

```

```

Training and evaluating SVM classifier with threshold 0.5
Confusion Matrix (Threshold = 0.5):
[[45 10]
 [14 41]]

```

```

Classification Report (Threshold = 0.5):
              precision    recall  f1-score   support

     0       0.76      0.82      0.79        55
     1       0.80      0.75      0.77        55

 accuracy          0.78
macro avg          0.78      0.78      0.78        110
weighted avg       0.78      0.78      0.78        110

```

```

Training and evaluating SVM classifier with threshold 0.6

```

Confusion Matrix (Threshold = 0.6):

```
[[47  8]
 [17 38]]
```

Classification Report (Threshold = 0.6):

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

```
for classifier, name in zip(classifiers, classifier_names):
    classifier.fit(X_train2, Y_train2)
```

```
# Plot ROC curves for all classifiers
```


```
plt.figure(figsize=(10, 6))
```

```
for classifier, name in zip(classifiers, classifier_names):
```

```
    plot_roc_curve(classifier, X_test2, Y_test2, name)
```

```
plt.legend(loc='lower right')
```

```
plt.show()
```



✓ Ablation Study for 80:20 and threshold=0.5

```
from sklearn.metrics import accuracy_score
def ablation_study(classifier, X_train, Y_train, X_test, Y_test):
    feature_accuracies = {}

    # Iterate over each feature
    for feature in X_train.columns:
        # Train the classifier with only the current feature
        classifier.fit(X_train[[feature]], Y_train)

        # Predict on the test set
        Y_pred = classifier.predict(X_test[[feature]])

        # Calculate accuracy
        accuracy = accuracy_score(Y_test, Y_pred)

        # Store accuracy for the current feature
        feature_accuracies[feature] = accuracy

    # Sort features by accuracy
    sorted_features = sorted(feature_accuracies, key=feature_accuracies.get, reverse=True)

    return sorted_features
```

```
# Convert arrays to DataFrames
```

```
X_train_df = pd.DataFrame(X_train, columns=df.columns[:-1])
```

```
X_test_df = pd.DataFrame(X_test, columns=df.columns[:-1])
```

```
# Perform ablation study for each classifier
for classifier, name in zip(classifiers, classifier_names):
    print("\nAblation study for", name, "classifier:")
    sorted_features = ablation_study(classifier, X_train_df, Y_train, X_test_df, Y_test)

    # Print ranked features
    print("Ranked features based on accuracy:")
    for i, feature in enumerate(sorted_features):
        print(f"{i+1}. {feature}")
```

Ablation study for KNN classifier:

Ranked features based on accuracy:

1. slp
2. cp
3. restecg
4. thalachh
5. age
6. fbs
7. exng
8. caa
9. thall
10. trtbps
11. oldpeak
12. chol
13. sex

Ablation study for SVM classifier:

Ranked features based on accuracy:

1. cp
2. thall
3. exng
4. oldpeak
5. caa
6. thalachh
7. slp
8. age
9. sex
10. restecg
11. trtbps
12. chol
13. fbs

Ablation study for Naive Bayes classifier:

Ranked features based on accuracy:

1. cp
2. thall
3. exng
4. oldpeak
5. thalachh
6. slp
7. caa
8. age
9. sex
10. restecg
11. trtbps
12. chol
13. fbs

Ablation study for Logistic Regression classifier:

Ranked features based on accuracy:

1. cp
2. thall
3. exng
4. oldpeak
5. caa
6. thalachh
7. age

✓ Bar Graph & Performance metrics values for all case.

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
```

```
# Define function to calculate performance metrics
```

```
def calculate_metrics(classifier, X_train, Y_train, X_test, Y_test, threshold):
```

```
    # Train the classifier
```

```
    classifier.fit(X_train, Y_train)
```

```
    # Predict probabilities
```

```
    Y_prob = classifier.predict_proba(X_test)[: , 1]
```

```
    # Apply threshold
```

```

Y_pred = (Y_prob > threshold).astype(int)

# Calculate metrics
accuracy = accuracy_score(Y_test, Y_pred)
precision = precision_score(Y_test, Y_pred)
recall = recall_score(Y_test, Y_pred)
f1 = f1_score(Y_test, Y_pred)
roc_auc = roc_auc_score(Y_test, Y_prob)

return accuracy, precision, recall, f1, roc_auc

```

```

# Initializing dictionaries to store metrics
metrics_80_20 = {'Accuracy': [], 'Precision': [], 'Recall': [], 'F1 Score': [], 'ROC-AUC': []}
metrics_70_30 = {'Accuracy': [], 'Precision': [], 'Recall': [], 'F1 Score': [], 'ROC-AUC': []}

# Loop over classifiers and threshold values
for classifier, name in zip(classifiers, classifier_names):
    for threshold in thresholds:
        # Calculate metrics for 80:20 split
        accuracy_80_20, precision_80_20, recall_80_20, f1_80_20, roc_auc_80_20 = calculate_metrics(classifier, X_train, Y_train, threshold)

        # Append metrics to dictionary for 80:20 split
        metrics_80_20['Accuracy'].append(accuracy_80_20)
        metrics_80_20['Precision'].append(precision_80_20)
        metrics_80_20['Recall'].append(recall_80_20)
        metrics_80_20['F1 Score'].append(f1_80_20)
        metrics_80_20['ROC-AUC'].append(roc_auc_80_20)

        # Calculate metrics for 70:30 split
        accuracy_70_30, precision_70_30, recall_70_30, f1_70_30, roc_auc_70_30 = calculate_metrics(classifier, X_train2, Y_train2, threshold)

        # Append metrics to dictionary for 70:30 split
        metrics_70_30['Accuracy'].append(accuracy_70_30)
        metrics_70_30['Precision'].append(precision_70_30)
        metrics_70_30['Recall'].append(recall_70_30)
        metrics_70_30['F1 Score'].append(f1_70_30)
        metrics_70_30['ROC-AUC'].append(roc_auc_70_30)

```

```

def plot_bar_graph(metrics_df, title):
    fig, ax = plt.subplots(figsize=(12, 6))

    # Set the width of each bar
    bar_width = 0.2
    index = np.arange(len(metrics_df))

    # Define color palette for bars
    colors = ['b', 'g', 'r', 'c', 'm']

    # Iterate over metrics and plot bars for each metric
    for i, (metric_name, metric_values) in enumerate(metrics_df.items()):
        # Calculate the position of bars for each classifier
        bar_positions = index + i * bar_width - (len(metrics_df.columns) - 1) * bar_width / 2

        # Plot bars for each classifier
        ax.bar(bar_positions, metric_values, bar_width, label=metric_name, color=colors[i])

    # Set labels and title
    ax.set_xlabel('Classifier')
    ax.set_ylabel('Score')
    ax.set_title(title)
    ax.set_xticks(index)
    ax.set_xticklabels(metrics_df.index)
    ax.legend()

    # Show plot
    plt.tight_layout()
    plt.show()

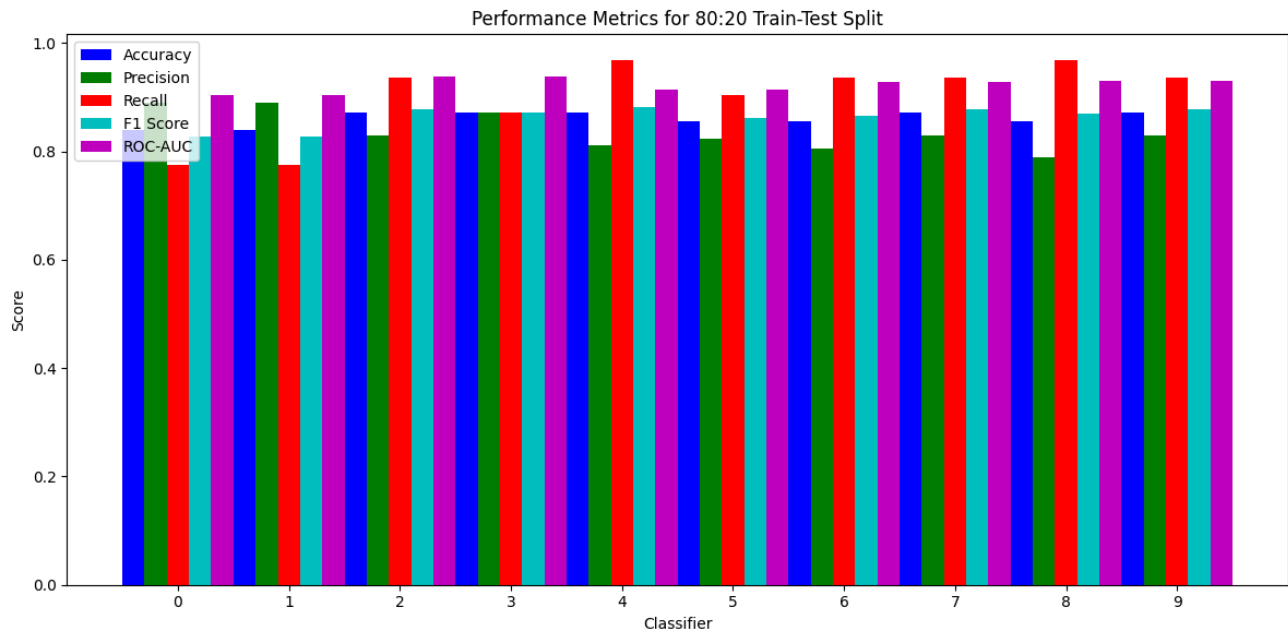
```

```


metrics_80_20_df = pd.DataFrame.from_dict(metrics_80_20, orient='index') # converting from dictionary to dataframe
metrics_70_30_df = pd.DataFrame.from_dict(metrics_70_30, orient='index')

```

plot_bar_graph(metrics_80_20_df.T, 'Performance Metrics for 80:20 Train-Test Split')



plot_bar_graph(metrics_70_30_df.T, 'Performance Metrics for 70:30 Train-Test Split')



print(metrics_80_20_df.T)

	Accuracy	Precision	Recall	F1 Score	ROC-AUC
0	0.838710	0.888889	0.774194	0.827586	0.904787
1	0.838710	0.888889	0.774194	0.827586	0.904787
2	0.870968	0.828571	0.935484	0.878788	0.938606
3	0.870968	0.870968	0.870968	0.870968	0.938606
4	0.870968	0.810811	0.967742	0.882353	0.913632
5	0.854839	0.823529	0.903226	0.861538	0.913632
6	0.854839	0.805556	0.935484	0.865672	0.928200
7	0.870968	0.828571	0.935484	0.878788	0.928200
8	0.854839	0.789474	0.967742	0.869565	0.929240
9	0.870968	0.828571	0.935484	0.878788	0.929240

